

# CS3236 Lecture Notes #2: Symbol-Wise Source Coding

Jonathan Scarlett

February 8, 2024

## Useful references:

- Cover/Thomas Chapter 5
- MacKay Chapter 5

## 1 Symbol-Wise Coding

### Setup.

- Consider a discrete random variable  $X$  with probability mass function (PMF)  $P_X$ .
  - For example, for text (without spaces/punctuation)  $X$  might take one of 26 characters  $\{a, \dots, z\}$ , with  $P_X(e)$  being highest,  $P_X(q)$  being low, etc.
  - More generally, the set of all symbols is denoted by  $\mathcal{X}$  (called the “alphabet” even when not referring to the English alphabet or even text).
- Symbol-wise source coding maps each  $x \in \mathcal{X}$  to some binary sequence  $C(x)$ . The length of this sequence is denoted by  $\ell(x)$ .
  - e.g., Map ‘x’ to 0011010 and ‘q’ to 0010100 because they are uncommon symbols, map ‘e’ to 1 because it is a very common symbol.

Since having a small length is so fundamental, we provide the following formal definition.

- **Definition.** The average length of a code  $C(\cdot)$  is given by

$$L(C) = \sum_{x \in \mathcal{X}} P_X(x) \ell(x).$$

- We need to be able to map back from the binary sequences to the original alphabet, so we cannot make every binary sequence short!
- Let’s look at the decoding conditions we would like to have.

## Decodability conditions.

- To have any hope of mapping the binary sequences back to the original alphabet, we need that  $C(x) \neq C(x')$  whenever  $x \neq x'$ . This condition is so trivial that it doesn't really need a name, but sometimes it's called the *nonsingular* property.
- We are actually interested in coding multiple alphabet symbols in succession, so being nonsingular is not enough. Consider the following code for  $\mathcal{X} = \{1, 2, 3, 4\}$ :

$$\begin{aligned} a &\rightarrow 0 \\ b &\rightarrow 1 \\ c &\rightarrow 00 \\ d &\rightarrow 11 \end{aligned}$$

Clearly there is no way to distinguish the sequence 'aabb' from 'cd'.

- **Definition.** A code  $C(\cdot)$  is said to be *uniquely decodable* if no two sequences (of equal or differing lengths) of symbols in  $\mathcal{X}$  are coded to the same concatenated binary sequence. That is,  $x_1, \dots, x_n$  can always uniquely be identified from the string  $C(x_1) \dots C(x_n)$ .
- **Example.** The following code is uniquely decodable:

$$\begin{aligned} a &\rightarrow 1 \\ b &\rightarrow 10 \\ c &\rightarrow 100 \\ d &\rightarrow 1000 \end{aligned}$$

While unique decodability is easy to see in this example, it can be tricky to verify in larger codes. It is more convenient to work with the following *seemingly* more restrictive condition.

- **Definition.** A code  $C(\cdot)$  is said to be *prefix-free* if no codeword is a prefix of any other (i.e., it is not possible to append more bits to some  $C(x)$  in order to produce some other  $C(x')$ ).
  - Sometimes the terminology *instantaneous code* is used to mean the same thing.
  - It turns out (we will omit the proof) that restricting to prefix-free codes instead of general uniquely decodable codes does not lose us anything in the average length we can achieve.
  - However, while decoding uniquely decodable codes is challenging in general, decoding prefix-free codes is trivial: As soon as the binary sequence matches some  $C(x)$ , output  $x$ , and then iteratively continue with the rest of the binary sequence.
  - We will therefore focus primarily on prefix-free codes.
- **Example.** The following code is prefix-free:

$$\begin{aligned} a &\rightarrow 0 \\ b &\rightarrow 10 \\ c &\rightarrow 110 \\ d &\rightarrow 111 \end{aligned}$$

Here is a simple example of decoding an encoded sequence for this code:

$cab \rightarrow \underline{110} \underline{010}$   
 $abba \rightarrow \underline{010} \underline{1010}$   
 $dad \rightarrow \underline{1110} \underline{111}$

- **Note:** We will focus on binary codes (which are by far the most common), but the vast majority of what we will cover can be easily extended to  $D$ -ary codes for  $D > 2$  (e.g., for  $D = 3$ , we have ternary codes with digits  $\{0, 1, 2\}$ , so we can have mappings like  $b \rightarrow 21$ ,  $q \rightarrow 0121$ ,  $z \rightarrow 2222$ ).

## 2 Kraft's Inequality and the Entropy Bound


### Kraft's Inequality.

- Clearly not all possible combinations of lengths are possible, e.g., we cannot map every letter  $a \dots z$  to a sequence of length 3 or less – there are not enough such sequences. Even when there are enough sequences, not all allocations of symbols to sequences will be prefix-free (or uniquely decodable). Kraft's inequality gives a useful condition that any prefix-free code must satisfy.
- **Theorem (Kraft's inequality).** Any prefix-free code  $C(\cdot)$  that maps each  $x \in \mathcal{X}$  to a codeword of length  $\ell(x)$  must satisfy

$$\sum_{x \in \mathcal{X}} 2^{-\ell(x)} \leq 1.$$

- **Proof:**

- Represent the codewords by a binary tree as follows:



- By the prefix-tree assumption, if there is a codeword at some point in the tree, there are no codewords further down the tree.
- Now, consider starting at the root and then repeatedly branching either way with probability  $\frac{1}{2}$  each until a codeword is hit.
- Clearly, the probability of a given length- $\ell(x)$  codeword being hit is exactly  $2^{-\ell(x)}$ .

- But the total probability of hitting codewords cannot exceed one, so summing up all the probabilities gives  $\sum_{x \in \mathcal{X}} 2^{-\ell(x)} \leq 1$ .

• **Theorem (Existence property).** If a given set of integers  $\{\ell(x)\}_{x \in \mathcal{X}}$  satisfies  $\sum_{x \in \mathcal{X}} 2^{-\ell(x)} \leq 1$ , then it is possible to construct a prefix-free code that maps each  $x \in \mathcal{X}$  to a codeword of length  $\ell(x)$ .

- Proof outline: Essentially proved by choosing codewords of a suitable length on a tree like the one shown above, starting with those having the smallest length. When  $\sum_{x \in \mathcal{X}} 2^{-\ell(x)} \leq 1$ , we never “run out of space” on the tree; see the relevant tutorial question for details.
- Hence, the condition in Kraft’s inequality is both necessary and sufficient for the existence of a prefix-free code having such lengths.

### Entropy bound.

• **Theorem.** For  $X \sim P_X$  and any prefix-free code  $C(\cdot)$ , the expected length satisfies

$$L(C) \geq H(X),$$

with equality if and only if  $P_X(x) = 2^{-\ell(x)}$  for all  $x \in \mathcal{X}$ .

- Hence, entropy provides a fundamental limit – we can never get an average length smaller than the entropy using a prefix-free code.
- Even though we are only stating/proving it for the prefix-free case, it can be shown that the same holds for any uniquely decodable code.

• **Proof.**

- Recall the definition of KL divergence,  $D(P\|Q) = \sum_x P(x) \log_2 \frac{P(x)}{Q_X(x)}$ .
- Observe that

$$\begin{aligned} L(C) - H(X) &\stackrel{(a)}{=} \sum_x P_X(x) \ell(x) - \sum_x P_X(x) \log_2 \frac{1}{P_X(x)} \\ &\stackrel{(b)}{=} \sum_x P_X(x) \log_2 2^{\ell(x)} - \sum_x P_X(x) \log_2 \frac{1}{P_X(x)}, \end{aligned}$$

where (a) is by definition, and (b) simply uses  $c = \log_2(2^c)$ .

- To simplify notation, let  $Z = \sum_{x \in \mathcal{X}} 2^{-\ell(x)}$ , and define  $Q_X(x) = \frac{2^{-\ell(x)}}{Z}$  which is a valid PMF (i.e., has non-negative values summing to one).
- Re-arranging terms in the definition of  $Q_X(x)$  gives  $2^{\ell(x)} = \frac{1}{Z \cdot Q_X(x)}$ , and substitution into the above equation gives

$$\begin{aligned} L(C) - H(X) &= \sum_x P_X(x) \log_2 \frac{1}{Z Q_X(x)} - \sum_x P_X(x) \log_2 \frac{1}{P_X(x)} \\ &\stackrel{(a)}{=} \log_2 \frac{1}{Z} + \sum_x P_X(x) \log_2 \frac{P_X(x)}{Q_X(x)} \\ &\stackrel{(b)}{=} \log_2 \frac{1}{Z} + D(P_X\|Q) \\ &\stackrel{(c)}{\geq} 0, \end{aligned}$$

where (a) uses simple re-arranging (and  $\log_2 \frac{1}{\alpha} = -\log_2 \alpha$ ), (b) uses the definition of KL divergence, and (c) uses  $Z \leq 1$  (by Kraft's inequality) and  $D(P_X \| Q_X) \geq 0$  (KL divergence between two PMFs is always non-negative).

- To get  $L(C) = H(X)$ , we need both  $Z \leq 1$  and  $D(P_X \| Q_X) \geq 0$  to hold with equality. The former condition gives  $Q_X(x) = 2^{-\ell(x)}$  (see the definition of  $Q_X$ ), and the latter gives  $P_X = Q_X$  (as established in the previous lecture), so overall we require  $P_X(x) = 2^{-\ell(x)}$  for all  $x$ .

- **Implication.** If our probabilities all contain powers of two ( $\frac{1}{2}, \frac{1}{4}, \frac{1}{8}, \frac{1}{16}$ , etc.), we can bring the average code length all the way down to the entropy.

- A simple example:

SYMBOL	PROBABILITY	CODEWORD
a	$\frac{1}{2}$	0
b	$\frac{1}{4}$	10
c	$\frac{1}{8}$	110
d	$\frac{1}{16}$	1110
e	$\frac{1}{16}$	1111

- Notice that the lengths satisfy  $\ell(x) = \log_2 \frac{1}{P_X(x)}$  for all  $x \in \{a, b, c, d, e\}$

### 3 Shannon-Fano Code

- Based on the “...equality if and only if...” statement in the entropy bound theorem, we can think of  $\ell^*(x) = \log_2 \frac{1}{P_X(x)}$  as being the “ideal” code length. However, it can only be attained when all values of  $\frac{1}{P_X(x)}$  are powers of two.
- The **Shannon-Fano code** simply rounds the ideal lengths up to the nearest integer:

$$\ell(x) = \left\lceil \log_2 \frac{1}{P_X(x)} \right\rceil,$$

where  $\lceil \cdot \rceil$  is the ceiling operation (i.e., rounding up).

- These lengths satisfy the conditions of the “Existence property” theorem above, since

$$\sum_{x \in \mathcal{X}} 2^{-\ell(x)} = \sum_{x \in \mathcal{X}} 2^{-\lceil \log_2 \frac{1}{P_X(x)} \rceil} \leq \sum_{x \in \mathcal{X}} 2^{-\log_2 \frac{1}{P_X(x)}} = \sum_{x \in \mathcal{X}} P_X(x) = 1,$$

where we used  $\lceil \alpha \rceil \geq \alpha$  and  $-\log_2 \frac{1}{\alpha} = \log_2 \alpha$ . Hence, that theorem implies that we can indeed construct a prefix-free code with the above lengths.

- **Theorem.** The average length  $L(C)$  of the Shannon-Fano code satisfies

$$H(X) \leq L(C) < H(X) + 1,$$

so is within one bit of the best average length possible.

- **Proof:** The lower bound is just a repetition of the entropy bound. To prove the upper bound, we use the fact that  $\lceil \alpha \rceil < \alpha + 1$  to deduce the following:

$$\begin{aligned} L(C) &= \sum_{x \in \mathcal{X}} P_X(x) \ell(x) \\ &= \sum_{x \in \mathcal{X}} P_X(x) \left\lceil \log_2 \frac{1}{P_X(x)} \right\rceil \\ &< \sum_{x \in \mathcal{X}} P_X(x) \left( \log_2 \frac{1}{P_X(x)} + 1 \right) \\ &= H(X) + 1, \end{aligned}$$

where the last step uses the definition of entropy and  $\sum_{x \in \mathcal{X}} P_X(x) = 1$ . (Note: By following similar steps but instead applying  $\lceil \alpha \rceil \geq \alpha$ , we get an alternative proof of the lower bound.)

- While the addition of at most 1 bit may seem innocuous, it can be very significant (e.g., for a “low-information” source, maybe  $H(X)$  itself is only 0.5 bits!)
- **Theorem (Mismatched case).** If the true distribution is  $P_X$  but the lengths are chosen according to  $Q_X$  (i.e.,  $\ell(x) = \lceil \log_2 \frac{1}{Q_X(x)} \rceil$ ), then the Shannon-Fano code satisfies

$$H(X) + D(P_X \| Q_X) \leq L(C) < H(X) + D(P_X \| Q_X) + 1.$$

- **Proof:** Similar to above, also using  $\mathbb{E}_P \left[ \log_2 \frac{1}{Q_X(X)} \right] = \mathbb{E}_P \left[ \log_2 \frac{P_X(X)}{Q_X(X)P_X(X)} \right] = \mathbb{E}_P \left[ \log_2 \frac{1}{P_X(X)} + \log_2 \frac{P_X(X)}{Q_X(X)} \right] = H(X) + D(P_X \| Q_X)$ .
- Hence, if an inaccurate distribution is used (not-so-small  $D(P_X \| Q_X)$ ) then the penalty due to mismatch may also be significant.

## 4 Huffman Code

- At this stage it is natural to ask whether it is possible to find the *optimal* symbol code, in the sense of minimizing  $L(C)$  while being uniquely decodable. This remained a seemingly challenging open problem until an *extremely simple* solution was given by Huffman (as part of a homework question!).
- **Huffman code.** Construct a tree as follows:
  - List the symbols of  $\mathcal{X}$  from highest probability from highest to lowest.
  - Draw a branch connecting the two symbols with the lowest probability, and label the merged point with the sum of the two associated probabilities.
  - Repeat the previous step (with the two original probabilities replaced by the merged probability) until everything has merged to a single point with total probability 1.

Once this tree is constructed, we label the two edges in each branch as 0 and 1, and then let the codewords be the labels encountered when traversing from the end back to the start.

- INIT:

0:6  
0:14  
0:1  
0:1  
0:06

(INIT)

STEP 1:

0:6  
0:14  
0:2  
0:16

(STEP 1)

STEP 2:

0:6  
0:4  
0:16

(STEP 2)

STEP 3:

0:6  
0:4

(STEP 3)

FINAL:

0:6  
0:4  
0:16  
0:2  
0:1  
0:1  
0:06  
1:0

(FINAL)

- ## Discussion

- $$NH(X) \leq L(C) \leq NH(X) + 1,$$

$$H(X) \leq \text{Average length per symbol} \leq H(X) + \frac{1}{N},$$


- Disadvantage 1: Determining the distribution  $P_{X_1, \dots, X_N}$  accurately is very difficult.

- Disadvantage 2: Even if the joint distribution is known, sorting  $|\mathcal{X}|^N$  probabilities in the Huffman coding algorithm becomes computationally challenging even for moderate values of  $N$ . The complexity increases exponentially with  $N$ .

## 6 (Optional) Beyond Symbol-Wise Codes

### Arithmetic codes

- Arithmetic codes are a very elegant technique for sequentially coding sources with memory when the conditional distribution  $P_{X_i|X_1, \dots, X_{i-1}}$  is known. The idea is illustrated in the following:



- For concreteness, suppose the alphabet is ‘a’, ‘b’, ‘c’.
- Start with an interval ranging from 0 to 1.
- Split the interval into three regions of width  $P_{X_1}(a)$ ,  $P_{X_1}(b)$ , and  $P_{X_1}(c)$ .
- After observing  $X_1 = b$ , move into the region of width  $P_{X_1}(b)$ .
- Split the width- $P_{X_1}(b)$  region into three regions proportional to  $P_{X_2|X_1}(a)$ ,  $P_{X_2|X_1}(b)$ , and  $P_{X_2|X_1}(c)$ .
- After observing  $X_2 = c$ , move into the corresponding sub-region.
- Continue recursively until the entire input sequence has been read.
- At the end of this process, we are left with a very small sub-interval  $\mathcal{I}$  of  $[0, 1]$ . How do we map this to a binary sequence?
- **Key idea.**
  - Every point in the interval  $[0, 1]$  corresponds to an infinite binary sequence (e.g.,  $\frac{1}{3}$  maps to  $0.010101\dots$ ,  $\frac{1}{2}$  maps to  $0.10000\dots$ , etc.).
  - A finite-length binary sequence (e.g.,  $0.010101$ ) then corresponds to an *interval* (e.g. starting from the number represented by  $0.010101$  followed by infinitely many zeros, ending at the number  $0.010101$  followed by infinitely many ones).



- Output a finite-length binary sequence that is just long enough for its corresponding interval to be a sub-interval of  $\mathcal{I}$ .
- The notion of entropy as a fundamental compression limit can be extended to broad types of sources with memory (see Cover/Thomas Chapter 4), and it can be shown that arithmetic coding can encode a sequence  $(x_1, \dots, x_n)$  down to at most

$$\ell(x_1, \dots, x_n) \leq \log_2 \frac{1}{P_{X_1, \dots, X_n}(x_1, \dots, x_n)} + 2$$

bits, which is within 2 bits of the “ideal length”. In particular, we get


$$\text{Average Total Length} \leq H(X_1, \dots, X_n) + 2,$$

and for a memoryless source, the number of bits per symbol is at most  $H(X) + \frac{2}{n}$ .

- See Cover/Thomas Section 13.3 or MacKay Section 6.2 for further details.

### Lempel-Ziv code

- A disadvantage of arithmetic codes is the need to know  $P_{X_1, \dots, X_n}$ . A class of codes known as *Lempel-Ziv* (LZ) codes are *universal*, in that they do not use any knowledge of the source distribution.
- For memoryless sources, and also broad classes of sources with memory, LZ codes are efficient enough to code almost down to the entropy (albeit with the “second-order” term being  $O(\log n)$ , which is a fair bit higher than the 2 bits attained by arithmetic coding).
- The encoding can roughly be described as follows:
  - Step 1: Parse the string into substrings that haven’t been observed earlier:



- Step 2: Encode each parsed sub-string into a sequence of bits of the form (pointer, new bit), where “pointer” identifies the index of the sub-string that matches the current one with the final bit removed, and “new bit” describes that final bit.
- Encoding: Store the sequence of pointers (represented in binary) and new bits.
- See Cover/Thomas Section 13.4 or MacKay Section 6.4 for further details.